

CAT-3

Portions


have to get
Byzantine consensus?

Failures


rollback (but send(m2) not in S2)

⇒ m2 is an orphaned msg

ie, only rcr recorded ⇒ inconsistent S2



Orphaned msg



system is consistent

lost msg

when rollback of S2 happens, it will expect rcr(m2), but it never happens.*

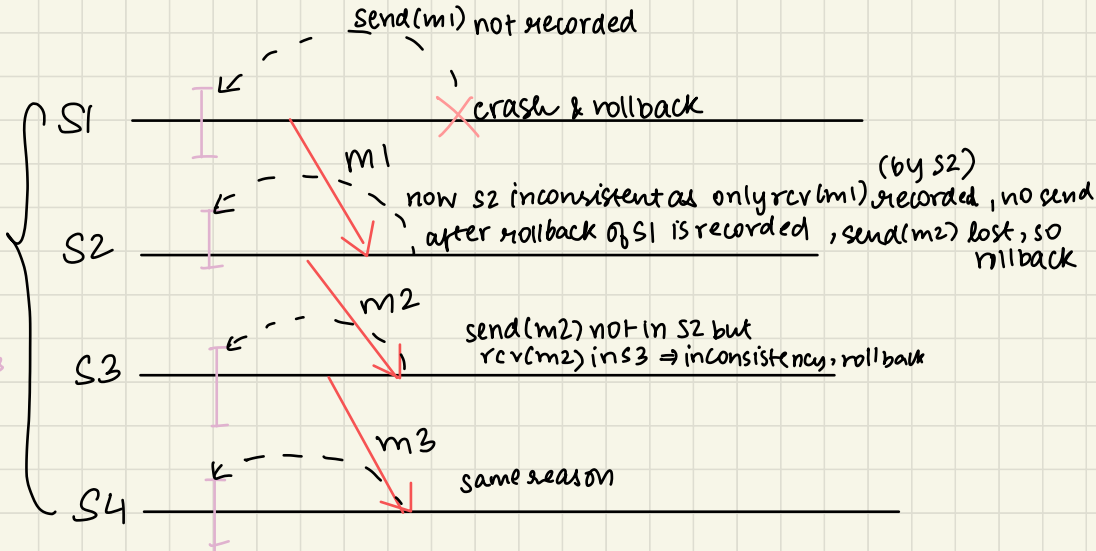
Doesn't affect consistency but affects performance

* ie, S2 assumes m1 in transit still

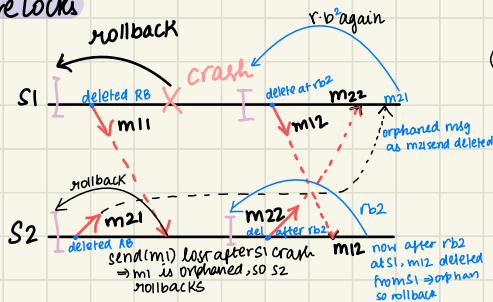
domino effect

aka

cascading rollbacks




Live locks



rollbacks keep reoccurring,

no progress (1 physical crash leads to many logical crashes)

Sync Checkpointing → only records useful cps (unlike MSRA)



GSRA → maintains consistency


but records unwarranted c.p.s

only
in GIRA

Sync Checkpointing - Algorithm



Initiator doesn't BC marker,
only sends through cohorts



(label of last msg \neq recd from x)

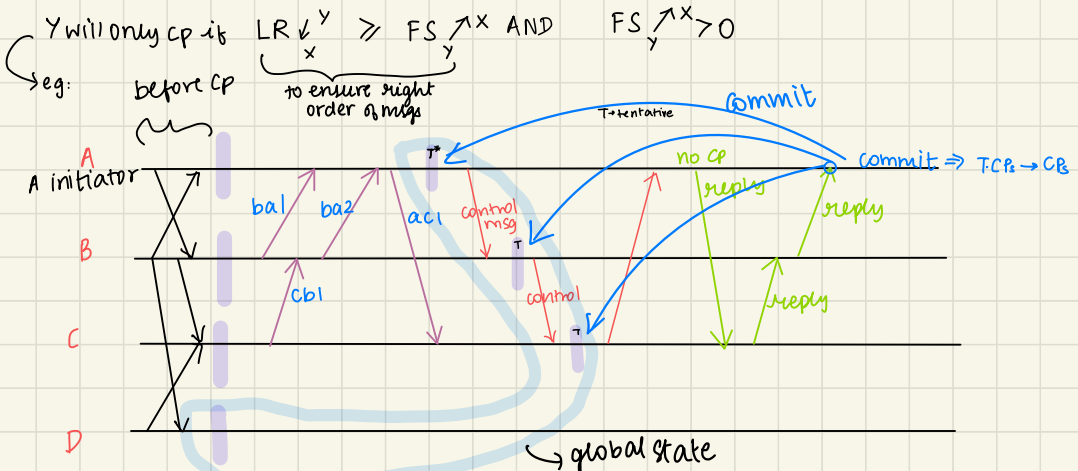
FS \uparrow y : 1 (label of first msg X sent to y)

$$FS \uparrow \begin{matrix} x \\ y \end{matrix} : 1$$

CASE - 1

Cohort_x: {Y | LR_x ↓ Y}

Y is a cohort of X
if X has rec'd a msg
from Y



after A cps, A sends M to its cohorts
here $\text{cohorts}_A = \{b\}$

So, M only sent to B from A

after B rcvs control msg,

check if condition holds,

check if condition holds,

LR $\downarrow^B \geq$ FS \uparrow^A AND FS \uparrow^A > 0
A B B

2 \geq 1 True 1 > 0 True

Since true, B cp.s & sends M to cohorts

$$\text{cohort}_B = \{C\}$$

after Crcvs controlmsg, check condition

$LR \downarrow^c_B \geq FS \uparrow^B_c$ AND $FS \uparrow^B_c > 0$
 $| \geq | \text{ true}$ $| > 0 \text{ true}$

$\Rightarrow C$ cps & send M to cohorts

Cohort_C = {A}

after A rcrs cp, check if

LR \downarrow^A_C \Rightarrow FS \nearrow^C_A AND FS $\nearrow^C_A > 0$

here, A already took CP, so no msg sent after that CP

$$\Rightarrow FS \downarrow_c^A = 0 \text{ after T. CP}$$

so $1 \geq 0$ true but $0 > 0$ false


A now doesn't take CP but replies to C to ack.

then C replies to B & B to A


after Initiator recvs a reply, it commits

after commit, tentative cp becomes

recorded as state



CASE 2



at A → initiator

$coh_A = \{B\}$

$LR_A^B = ba2 \rightarrow 2$ (msg B to inform)

at B

rcv msg from A

$LR_A^B \geq FS_B^A \ \&\& \ FS_B^A > 0$
 $2 \geq 1 \ \&\& \ 1 > 0$

checkpoint

$coh_B = \{C, A\}$

A already cp'd, but C hasn't

$LR_B^A = 1$ (send A)

$LR_B^C = 1$ (send C)

concurrently

at C

check

$LR_B^C \geq FS_C^B \ \&\& \ FS_C^B > 0$
 $1 \geq 1 \ \&\& \ 1 > 0 \checkmark$

checkpoint

$coh_C = \{\emptyset\}$

terminate & ack B

at A

check

$LR_B^A \geq FS_A^B \ \&\& \ FS_A^B > 0$

Only consider after tentative

cp of A $\Rightarrow FS_A^B$ after cp = 0

$1 > 0 \times$ no cp.

ack B & terminate

example



at A

$$\text{con}_A = \{D\}$$
$$LR_A^D = 1 \text{ (send to D)}$$

at D

$LR \downarrow^D \geq FS \uparrow^A$ $\&\& FS \uparrow^A > 0$

A D P

$17/1 \checkmark$ $170 \checkmark$

checkpoint

$$\text{coh } D = \{C, B\}$$
$$LR_D^C = 1 \quad LR_D^B = 1$$



→ send to C → send to B

Concurrently

at C

cut B

Sync Recovery



NOTE: alerts sent through all outgoing channels

$$LS_A^B \rightarrow ab1 = 1$$

$$LR_B^A \rightarrow ab1 = 1$$

after A's crash, B has to become aware of it. How?

↳ A can inform B after it recovers, (if B is A's cohort) → Case 1


↳ If A can't recover, (2nd case)

For B's recovery

① $LR_B^A > LS_A^B \Rightarrow B \text{ has to recover}$

last sent A to B after A's crash

$1 > 0 \Rightarrow B \text{ has to recover, i.e., rollback to cp. (send alert to C \& D)}$



↳ msg is lost after C's crash but this doesn't affect consistency, only performance

② a

② b concurrently rcv alert from B

at C, (from B)

$$LR_C^B > LS_B^C ?$$

$$1 > 0 \rightarrow \text{true}$$

⇒ C has to recover

↳ rollback (alert D)

③ at D (from C)

$$LR_D^C > LS_C^D ?$$

$$0 > 0 \rightarrow \text{false}$$

no rollback


at D (from B)

$$LR_D^B > LS_B^D ?$$

$$0 > 0 \rightarrow \text{false}$$

no rollback

example:



①a at B (alert from A)

$$LR \underset{B}{\swarrow}^A > LS \underset{A}{\nearrow}^B ?$$

0 > 0 → false

no rollback

①_b at C (alt from A)

$$LR \downarrow^A_C \quad 7LS \uparrow^C_A ?$$

1 > 0 ^{after crash} → true

rollback & alert B.

② at B (alert from C)

LR \downarrow^C 7 LS \uparrow^B ?
B C

1 70 → true

Rollback & alert D

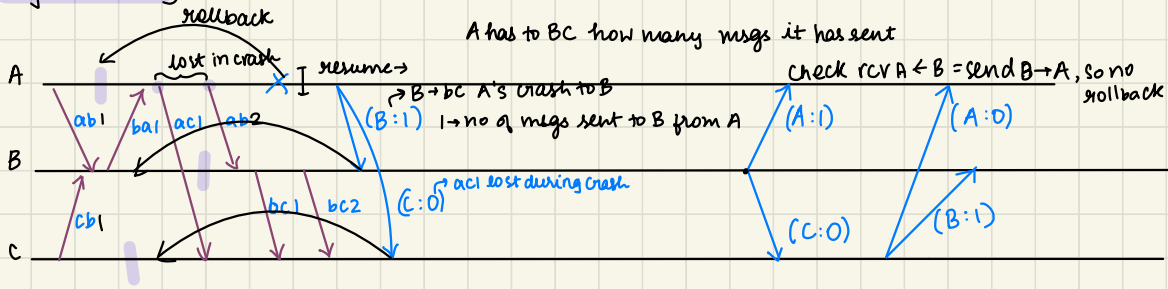
③ at D (alert from B)

$$LR \overset{B}{\nwarrow} \underset{D}{\nearrow} > LS \overset{D}{\nwarrow} \underset{B}{\nearrow} ?$$

$0 > 0 \rightarrow \text{false}$

no rollback

Async Recovery



at B,

send $A \rightarrow B = 1$

rcv $B \leftarrow A = 2$

send \leftarrow rcr \Rightarrow rollback (but to where?)

\hookrightarrow not to its prev cp as there is no coordination b/w checkpoints till any pt before rcr of msg $ab2$

at C

send $A \rightarrow C = 0$

rcv $C \leftarrow A = 1$

rcv \rightarrow send \Rightarrow rollback