

# Challenges: System Perspective (1)

- Communication mechanisms: E.g., Remote Procedure Call (RPC), remote object invocation (ROI), message-oriented vs. stream-oriented communication
- Processes: Code migration, process/thread management at clients and servers, design of software and mobile agents
- Naming: Easy to use identifiers needed to locate resources and processes transparently and scalably
- Synchronization
- Data storage and access
  - ▶ Schemes for data storage, search, and lookup should be fast and scalable across network
  - ▶ Revisit file system design
- Consistency and replication
  - ▶ Replication for fast access, scalability, avoid bottlenecks
  - ▶ Require consistency management among replicas

## Challenges: System Perspective (2)

- Fault-tolerance: correct and efficient operation despite link, node, process failures
- Distributed systems security
  - ▶ Secure channels, access control, key management (key generation and key distribution), authorization, secure group management
- Scalability and modularity of algorithms, data, services
- Some experimental systems: Globe, Globus, Grid

# Challenges: System Perspective (3)

- API for communications, services: ease of use
- Transparency: hiding implementation policies from user
  - ▶ Access: hide differences in data rep across systems, provide uniform operations to access resources
  - ▶ Location: locations of resources are transparent
  - ▶ Migration: relocate resources without renaming
  - ▶ Relocation: relocate resources as they are being accessed
  - ▶ Replication: hide replication from the users
  - ▶ Concurrency: mask the use of shared resources
  - ▶ Failure: reliable and fault-tolerant operation

# Challenges: Algorithm/Design (1)

- Useful execution models and frameworks: to reason with and design correct distributed programs
  - ▶ Interleaving model
  - ▶ Partial order model
  - ▶ Input/Output automata
  - ▶ Temporal Logic of Actions
- Dynamic distributed graph algorithms and routing algorithms
  - ▶ System topology: distributed graph, with only local neighborhood knowledge
  - ▶ Graph algorithms: building blocks for group communication, data dissemination, object location
  - ▶ Algorithms need to deal with dynamically changing graphs
  - ▶ Algorithm efficiency: also impacts resource consumption, latency, traffic, congestion

## Challenges: Algorithm/Design (2)

- Time and global state
  - ▶ 3D space, 1D time
  - ▶ Physical time (clock) accuracy
  - ▶ Logical time captures inter-process dependencies and tracks relative time progression
  - ▶ Global state observation: inherent distributed nature of system
  - ▶ Concurrency measures: concurrency depends on program logic, execution speeds within logical threads, communication speeds

# Challenges: Algorithm/Design (3)

- Synchronization/coordination mechanisms

- ▶ Physical clock synchronization: hardware drift needs correction
- ▶ Leader election: select a distinguished process, due to inherent symmetry
- ▶ Mutual exclusion: coordinate access to critical resources
- ▶ Distributed deadlock detection and resolution: need to observe global state; avoid duplicate detection, unnecessary aborts
- ▶ Termination detection: global state of quiescence; no CPU processing and no in-transit messages
- ▶ Garbage collection: Reclaim objects no longer pointed to by any process

# Challenges: Algorithm/Design (4)

- Group communication, multicast, and ordered message delivery
  - ▶ Group: processes sharing a context, collaborating
  - ▶ Multiple joins, leaves, fails
  - ▶ Concurrent sends: semantics of delivery order
- Monitoring distributed events and predicates
  - ▶ Predicate: condition on global system state
  - ▶ Debugging, environmental sensing, industrial process control, analyzing event streams
- Distributed program design and verification tools
- Debugging distributed programs

# Challenges: Algorithm/Design (5)

- Data replication, consistency models, and caching
  - ▶ Fast, scalable access;
  - ▶ coordinate replica updates;
  - ▶ optimize replica placement
- World Wide Web design: caching, searching, scheduling
  - ▶ Global scale distributed system; end-users
  - ▶ Read-intensive; prefetching over caching
  - ▶ Object search and navigation are resource-intensive
  - ▶ User-perceived latency



# Challenges: Algorithm/Design (6)

- Distributed shared memory abstraction
  - ▶ Wait-free algorithm design: process completes execution, irrespective of actions of other processes, i.e.,  $n - 1$  fault-resilience
  - ▶ Mutual exclusion
    - ★ Bakery algorithm, semaphores, based on atomic hardware primitives, fast algorithms when contention-free access
  - ▶ Register constructions
    - ★ Revisit assumptions about memory access
    - ★ What behavior under concurrent unrestricted access to memory?  
Foundation for future architectures, decoupled with technology (semiconductor, biocomputing, quantum ...)
  - ▶ Consistency models:
    - ★ coherence versus access cost trade-off
    - ★ Weaker models than strict consistency of uniprocessors

# Challenges: Algorithm/Design (7)

- Reliable and fault-tolerant distributed systems
  - ▶ Consensus algorithms: processes reach agreement in spite of faults (under various fault models)
  - ▶ Replication and replica management
  - ▶ Voting and quorum systems
  - ▶ Distributed databases, commit: ACID properties
  - ▶ Self-stabilizing systems: "illegal" system state changes to "legal" state; requires built-in redundancy
  - ▶ Checkpointing and recovery algorithms: roll back and restart from earlier "saved" state
  - ▶ Failure detectors:
    - ★ Difficult to distinguish a "slow" process/message from a failed process/ never sent message
    - ★ algorithms that "suspect" a process as having failed and converge on a determination of its up/down status

# Challenges: Algorithm/Design (8)

- Load balancing: to reduce latency, increase throughput, dynamically. E.g., server farms
  - ▶ Computation migration: relocate processes to redistribute workload
  - ▶ Data migration: move data, based on access patterns
  - ▶ Distributed scheduling: across processors
- Real-time scheduling: difficult without global view, network delays make task harder
- Performance modeling and analysis: Network latency to access resources must be reduced
  - ▶ Metrics: theoretical measures for algorithms, practical measures for systems
  - ▶ Measurement methodologies and tools