

Causal, Partial and Total Ordering of Messages in Distributed Systems

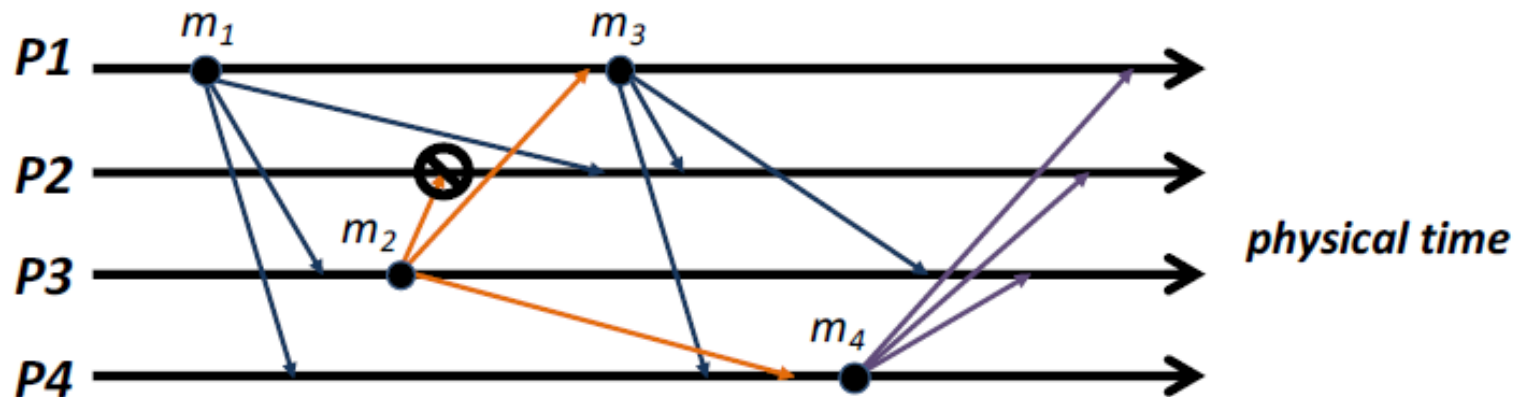
Causal and Total Order

Stronger orderings

- Can also implement FIFO ordering by just using a reliable FIFO transport like TCP/IP
- But the general ‘receive versus deliver’ model also allows us to provide **stronger** orderings:
 - **Causal ordering**: if event $\text{multicast}(g, m_1) \rightarrow \text{multicast}(g, m_2)$, then all processes will see m_1 before m_2
 - **Total ordering**: if any processes delivers a message m_1 before m_2 , then all processes will deliver m_1 before m_2
- Causal ordering implies FIFO ordering, since any two multicasts by the same process are related by \rightarrow
- Total ordering (as defined) does *not* imply FIFO (or causal) ordering, just says that all processes must agree
 - Often want **FIFO-total** ordering (combines the two)

Causal Ordering

Causal ordering




- Same example as previously, but now causal ordering means that
 - (a) everyone must see m_1 before m_3 (as with FIFO), **and**
 - (b) everyone must see m_1 before m_2 (due to happens-before)
- Is this ok?
 - No! $m_1 \rightarrow m_2$, but **P2** sees m_2 before m_1
 - To be correct, must hold back (delay) delivery of m_2 at **P2**
 - But how do we know this?

Causal Ordering

Implementing causal ordering

- Turns out this is pretty easy!
 - Start with receive algorithm for FIFO multicast...
 - and replace sequence numbers with vector clocks



- Some care needed with dynamic groups

Partial Order

partial order:

$\{A, B, C, D, E, F, G, H\}$

→ is an irreflexive partial order

- A set S , together with
- A binary relation, often written \leq , that lets you compare elements of S , and has the following properties:

X
vacuously true ✓
✓
→

- Reflexivity: for all $a \in S$, $a \leq a$.

- Antisymmetry: for all $a, b \in S$,
if $a \leq b$, and $b \leq a$,
 $a = b$

- Transitivity: for all $a, b, c \in S$,
if $a \leq b$ and $b \leq c$,
then $a \leq c$.

Total Order

Total ordering of events

- A system of clocks that satisfy the Clock Condition can be used to totally order system events.
- To totally order the events in a system, the events are ordered according to their times of occurrence. In case two or more events occur at the same time, an arbitrary total ordering \prec of processes is used. To do this, the relation \Rightarrow is defined as follows:

If a is an event in process P_i and b is an event in process P_j , then $a \Rightarrow b$ if and only if either:

- i. $C_i\langle a \rangle < C_j\langle b \rangle$ or
- ii. $C_i\langle a \rangle = C_j\langle b \rangle$ and $P_i \prec P_j$

There is total ordering because for any two events in the system, it is clear which happened first.

- The total ordering of events is very useful for distributed system implementation.