


Synchronizing Physical Clock Logical Clock & Vector Clock

Y. V. Lokeswari

ASP/CSE



Overview

- Physical Clocks
 - Synchronizing Physical Clock (Algorithms)
 - Problems with Physical Clock
 - Lamport's Logical Clock
 - Problems with Logical Clock
 - Vector Clock
 - Drawbacks of Vector Clocks
 - Applications of Vector Clocks
- 


Clock Synchronization

- Temporal ordering of events produced by concurrent processes
- Synchronization between senders and receivers of messages
- Serialization of concurrent access for shared objects
- **Physical Clock:** It is the internal clock present in a computer.
(Time of a day)
- **Logical Clock:** keeps track of event ordering among related
(causal) events.



Clock Synchronization

- Getting two systems to agree on time
 - Two clocks hardly ever agree
 - **Quartz oscillators oscillate at slightly different frequencies**
- **Clock Drift** : Clocks tick at different rates. Create ever-widening gap in perceived time
 - Clock drift in ordinary clocks based on quartz crystal is 10^{-6} seconds.
 - This creates a difference of 1 sec for every 11.6 days (1,000,000 sec)
 - Clock drift of high precision clock is 10^{-7} to 10^{-8}
- **Clock Skew** : Difference between two clocks at one point in time.



Clock Synchronization


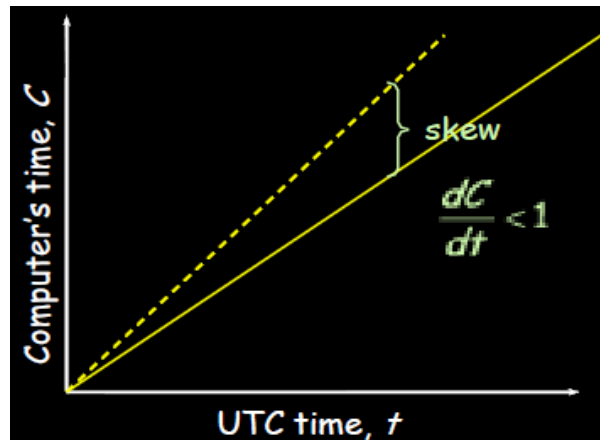



Figure 3.8 The behavior of fast, slow, and perfect clocks with respect to UTC.

Clock Synchronization


- Dealing with Clock Drift: Go for gradual clock correction

UTC : Universal Coordinated Time



If slow:

Make clock run faster until
it Synchronizes




If fast:

Make clock run slower
until it Synchronizes

Clock Synchronization

- **External Synchronization** : Clock synchronizes to correct time from external timing elements like Radio / Satellite time.
- **Internal Synchronization** : Clock synchronizes to correct time by getting timings from other computers.
- 3 - Algorithms for Synchronizing Physical Clock
 1. Cristian's Algorithm
 2. Berkeley Algorithm
 3. Network Time Protocol (NTP)




Clock Synchronization

1. Cristian's Algorithm

$(T_1 - T_0)/2$ is round-trip time

Client sets time to:



$$T_{new} = T_{server} + \frac{T_1 - T_0}{2}$$

Clock Synchronization

1. Cristian's Algorithm – Example

- Send request at 5:08:15.100 (T₀)
- Receive response at 5:08:15.900 (T₁)
- Response contains 5:09:25.300 (T_{server})
- Elapsed time is T₁ - T₀

$$5:08:15.900 - 5:08:15.100 = 800 \text{ msec}$$

- Best guess: timestamp was generated

$$400 \text{ msec ago } (800/2)$$

- Set time to T_{server} + elapsed time

$$5:09:25.300 + 400 = 5:09:25.700$$


Clock Synchronization


2. Berkeley Algorithm

- Machines run **time daemon** Process that implements protocol
- One machine is elected (or designated) as the server (**master**) Others are **slaves**
- Master polls each machine periodically and ask each machine for time
- Can use Cristian's algorithm to compensate for network latency
- When results are received, master computes average of times - Including master's time
- **Hope**: Average cancels out individual clock's tendencies to run fast or slow (clock drift)
- Send offset by which each clock needs adjustment to each slave.
- Algorithm has provisions for ignoring readings from clocks whose skew is too great
- Compute a **fault-tolerant average**
- If master fails - any slave can take over




Clock Synchronization

2. Berkeley Algorithm Example



Clock Synchronization

2. Berkeley Algorithm Example




2. Compute fault-tolerant average:

$$\frac{3:25 + 2:50 + 3:00}{3} = 3:05$$

Clock Synchronization

2. Berkeley Algorithm Example



3. Send offset to each client

Clock Synchronization

3. Network Time Protocol


- Enable clients across Internet to be accurately synchronized to UTC despite message delays
- Primary servers are connected directly to a time source such as a radio clock receiving UTC; secondary servers are synchronized with primary servers.
- The servers are connected in a logical hierarchy called a **synchronization subnet** whose levels are called **strata**.
- Primary servers occupy stratum 1: they are at the root.
- Stratum 2 servers are secondary servers that are synchronized directly with the primary servers;
- Stratum 3 servers are synchronized with stratum 2 servers, and so on.



Clock Synchronization

3. Network Time Protocol

Arrows denote synchronization control, numbers denote strata.



Clock Synchronization

Network Time Protocol

Figure 3.9 Offset and delay estimation [15].





Figure 3.10 Timing diagram for the two servers [15].



Let $a = T_1 - T_3$ and $b = T_2 - T_4$.

$$\theta = \frac{a+b}{2}, \quad \delta = |a-b|$$

the offset O_i can be estimated as

$$O_i = (T_{i-2} - T_{i-3} + T_{i-1} - T_i)/2.$$

The round-trip delay is estimated as

$$D_i = (T_i - T_{i-3}) - (T_{i-1} - T_{i-2}).$$

Clock Synchronization

3. Network Time Protocol (Synchronization Modes)

- **Multicast mode**
 - Every node sends its time to all the other nodes in the LAN
 - For high speed LANs
 - Lower accuracy but efficient
- **Procedure call mode**
 - Similar to Cristian's algorithm
- **Symmetric mode**
 - Intended for master servers
 - Pair of servers exchange messages and retain data to improve synchronization over time


All messages delivered unreliably with UDP



Problems with Physical Clocks

- **Clock Drift** : Clocks tick at different rates. Create ever-widening gap in perceived time
- **Clock Skew** : Difference between two clocks at one point in time.
- For quartz crystal clocks, typical drift rate is about one second every 10^6 seconds = 11.6 days
- Best atomic clocks have drift rate of one second in 10^{13} seconds = 300,000 years.
- **Quartz clock run at rate of 1.5 microseconds slower for every 35 days.**

Problems with Physical Clocks



8:01:24

Skew = +84 seconds
+84 seconds/35 days
Drift = +2.4 sec/day


Oct 23, 2006
8:00:00



8:01:48

Skew = +108 seconds
+108 seconds/35 days
Drift = +3.1 sec/day

Logical Clocks



Need for Logical Clock

- For many purposes, it is sufficient to know the order in which events occurred.
- Lamport (1978) — introduce logical (*virtual*) time, synchronize *logical clocks*.
- An event may be an instruction execution, may be a function execution, etc.
- Events include message send / receive

Within a single process, or between two processes on the same computer

- The order in which two **events** occur **can** be determined using the **physical clock**

Between two different computers in a distributed system

- The order in which two events occur **cannot** be determined **using local physical clocks**, since those clocks cannot be synchronized perfectly



Happened Before Relation

- Lamport defined the happened before relation (denoted as “ \rightarrow ”), which describes a **causal ordering of events**:
 1. if **a** and **b** are events in the same process, and **a** occurred before **b**, then **a** \rightarrow **b**
 2. if **a** is the event of sending a message **m** in one process, and **b** is the event of receiving that message **m** in another process, then **a** \rightarrow **b**
 3. if **a** \rightarrow **b**, and **b** \rightarrow **c**, then **a** \rightarrow **c** (i.e., the relation “ \rightarrow ” is transitive)



Causality

- Past events influence future events
- This influence among causally related events (those that can be ordered by “ \rightarrow ”) is referred to **causally affects**.
- If $a \rightarrow b$, event **a** causally affects event **b**


Concurrent events

- Two distinct events **a** and **b** are said to be concurrent (denoted “ $a \parallel b$ ”), if neither $a \rightarrow b$ nor $b \rightarrow a$
- In other words, concurrent events do not causally affect each other
- For any two events a and b in a system, either: $a \rightarrow b$ or $b \rightarrow a$ or $a \parallel b$



Causal Events

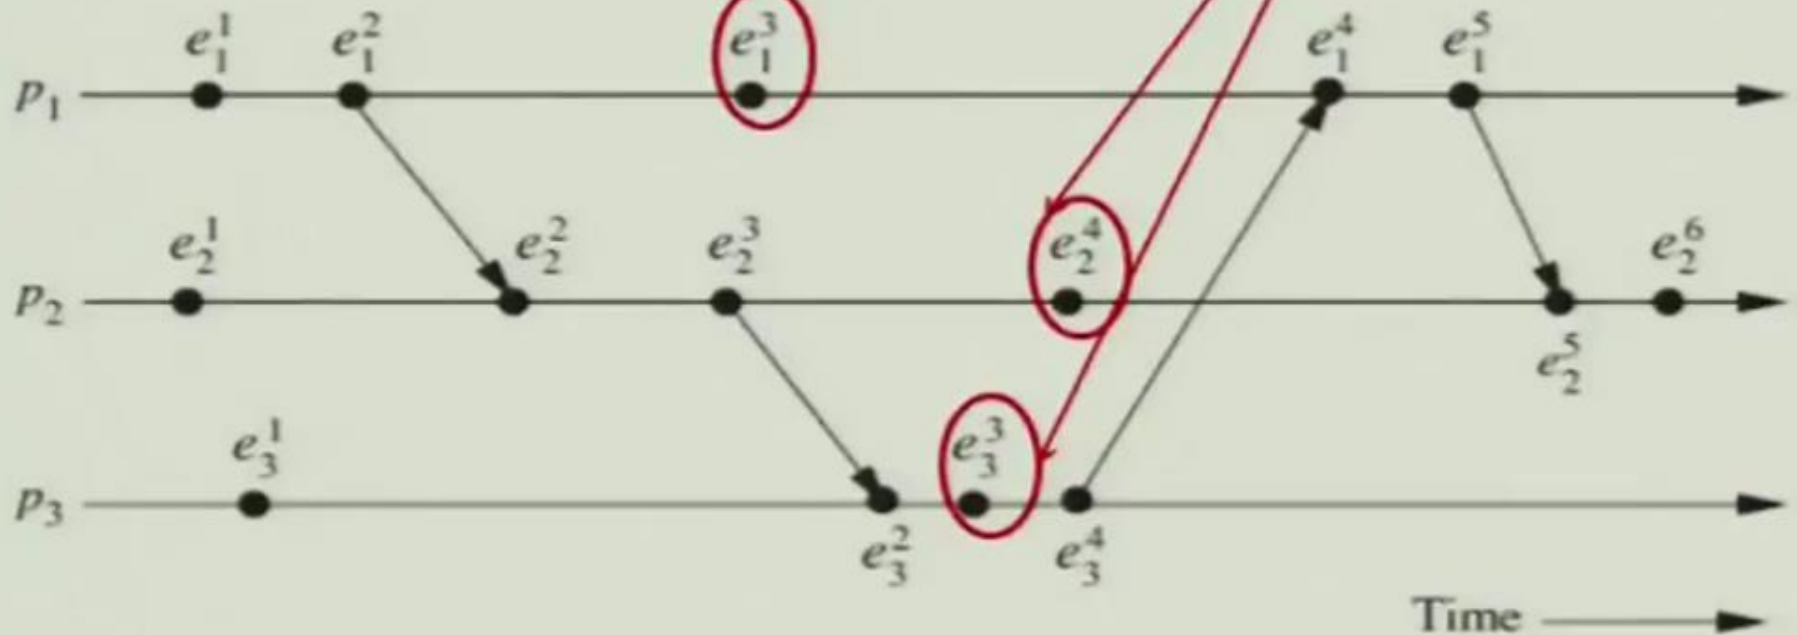
Space-time diagram



Concurrent Events

Space-time diagram

events in the set $\{e_1^3, e_2^4, e_3^3\}$ are logically concurrent.



Lamport's Logical Clock

- To implement “ \rightarrow ” in a distributed system, Lamport (1978) introduced the concept of logical clocks, which captures “ \rightarrow ” numerically
- Each process P_i has a logical clock C_i
- Clock C_i can assign a value $C_i(a)$ to any event **a** in process P_i
- The value $C_i(a)$ is called the timestamp of event **a** in process P_i
- The value $C(a)$ is called the timestamp of event **a** in whatever process it occurred.
- The timestamps have no relation to physical time, which leads to the term logical clock.
- The logical clocks assign monotonically increasing timestamps, and can be implemented by simple counters

Lamport's Logical Clock

- **Clock condition:** if $a \rightarrow b$, then $C(a) < C(b)$
- If event a happens before event b , then the clock value (timestamp) of a should be less than the clock value of b
- Note that we **can not say: if $C(a) < C(b)$, then $a \rightarrow b$**
- **Correctness conditions (must be satisfied by the logical clocks to meet the clock condition above):**
- [C1] For any two events a and b in the same process P_i ,
if a happens before b , then $C_i(a) < C_i(b)$
- [C2] If event a is the event of sending a message m in process P_i , and event b is the event of receiving that same message m in a different process P_k , then
 $C_i(a) < C_k(b)$

Lamport's Logical Clock

- **[IR1]** Clock C_i must be incremented between any two successive events in process P_i

$$C_i := C_i + d, \quad (d > 0) \text{ (usually } d=1)$$


- **[IR2]** If event a is the event of sending a message m in process P_i , then message m is assigned a timestamp $C_{msg} = C_i(a)$. When that same message m is received by a different process P_k , C_k is set to a value greater than or equal to its present value, and greater than C_{msg} .

$$C_k := \max(C_k, C_{msg} + d), \quad (d > 0) \text{ (usually } d=1)$$


Lamport's Logical Clock

IR 1: $C_i := C_i + d$, ($d > 0$), (usually $d=1$)


IR2: $C_k := \max(C_k, C_{msg} + d)$, ($d > 0$) (Usually $d=1$)



Lamport's Logical Clock

IR 1: $C_i := C_i + d$, ($d > 0$), (usually $d=1$)

IR2: $C_k := \max(C_k, C_{msg} + d)$, ($d > 0$) (Usually $d=1$)




Lamport's Logical Clock

- A total order of events (“ \Rightarrow ”) can be obtained as follows:
- If **a** is any event in process **P_i**, and **b** is any event in process **P_k**, then **a** \Rightarrow **b** if and only if either:
- $C_i(a) < C_k(b)$ or
- $C_i(a) = C_k(b)$ and $P_i \ll P_k$ (if scalar timestamps of **a** and **b** are equal, ordering of events need to done through some mechanism).

where “ \ll ” denotes a relation that totally orders the processes. Process identifier is used to break ties.

Lower the process identifier in ranking higher the priority.




Limitations of Lamport's Logical Clock

- With Lamport's logical clocks,
- if $a \rightarrow b$, then $C(a) < C(b)$
- The following is **not necessarily true if events a and b occur in** different processes:
- if $C(a) < C(b)$, then $a \rightarrow b$ **is NOT TRUE**, if **a** and **b** events occur in two different processes.
- Scalar time is not strongly consistent.




Limitations of Lamport's Logical Clock



- With Lamport's logical clocks, if $a \rightarrow b$, then $C(a) < C(b)$
 - ◆ The following is **not** necessarily true if events a and b occur in different processes: if $C(a) < C(b)$, then $a \rightarrow b$
 - ◆ $C(e_{11}) < C(e_{22})$, and $e_{11} \rightarrow e_{22}$ is true
 - ◆ $C(e_{11}) < C(e_{32})$, but $e_{11} \rightarrow e_{32}$ is false
- Cannot determine whether two events are causally related from timestamps

Limitations of Lamport's Logical Clock

There is **drastic increase in clock values** due to events that occur in different processes and that causally affects the other events. This is **not captured** in logical clocks.



Vector Clocks



Vector Clock

- Maintain a vector of values for every event that happens in all processes.
- Update happens for group of values in every event.



Vector Clock

- [IR1] Clock C_i must be incremented between any two successive events in process P_i :

$$C_i[i] := C_i[i] + d, \quad (d > 0, \text{ usually } d = 1)$$

- [IR2] If event a is the event of sending a message m in process P_i , then message m is assigned a vector timestamp $tm = C_i(a)$. When that same message m is received by a different process P_k , C_k is updated as follows:


$$\forall p, C_k[p] := \max(C_k[p], tm[p] + d), \quad (\text{usually } d = 0 \text{ unless needed to model network delay})$$

- It can be shown that $\forall i, \forall k : C_i[i] \geq C_k[i]$
- Rules for comparing timestamps can also be established so that
- if $ta < tb$, then $a \rightarrow b$ / Solves the problem with Lamport's clocks

Vector Clock

IR1: $C_i[i] := C_i[i] + d$. ($d=1$)

IR2: $\forall p, C_k[p] := \max(C_k[p], tm[p] + d)$. ($d=0$)




"enn" is
event;
"(n,n,n)" is
clock value

Vector Clock – Solving Logical Clock Problem

if $t(e_{31}) < t(e_{23})$ then $e_{31} \rightarrow e_{23}$ is true.

if $t(e_{12}) < t(e_{24})$ then $e_{12} \rightarrow e_{24}$ is also true.

Concurrent events $e_{32} \parallel e_{24}$; $e_{32} \not\rightarrow e_{24}$ and $e_{24} \not\rightarrow e_{32}$




"enn" is
event;
"(n,n,n)" is
clock value

Vector Clock – Solving Logical Clock Problem

if $C(e_{11}) < C(e_{22})$ then $e_{11} \rightarrow e_{22}$ is true.

if $C(e_{11}) < C(e_{32})$ then $e_{11} \rightarrow e_{32}$ is also true.

Vector time is strongly consistent.




Vector Clock – Concurrent Events

Concurrent events $e_{32} \parallel e_{24}$; $e_{32} \not\rightarrow e_{24}$ and $e_{24} \not\rightarrow e_{32}$

$C(e_{32}) = (0,0,2)$ and $C(e_{24}) = (2,4,1)$

$C(e_{32})$ is neither less nor greater than $C(e_{24})$

$e_{11} \parallel e_{21}$, $e_{11} \parallel e_{31}$, $e_{11} \parallel e_{32}$, $e_{21} \parallel e_{32}$, $e_{12} \parallel e_{32}$,
 $e_{22} \parallel e_{32}$, $e_{32} \parallel e_{13}$.



"enn" is
 event;
 "(n,n,n)" is
 clock value

Vector Clock

Example of Vector Clock




Figure 4.3: Evolution of vector time.

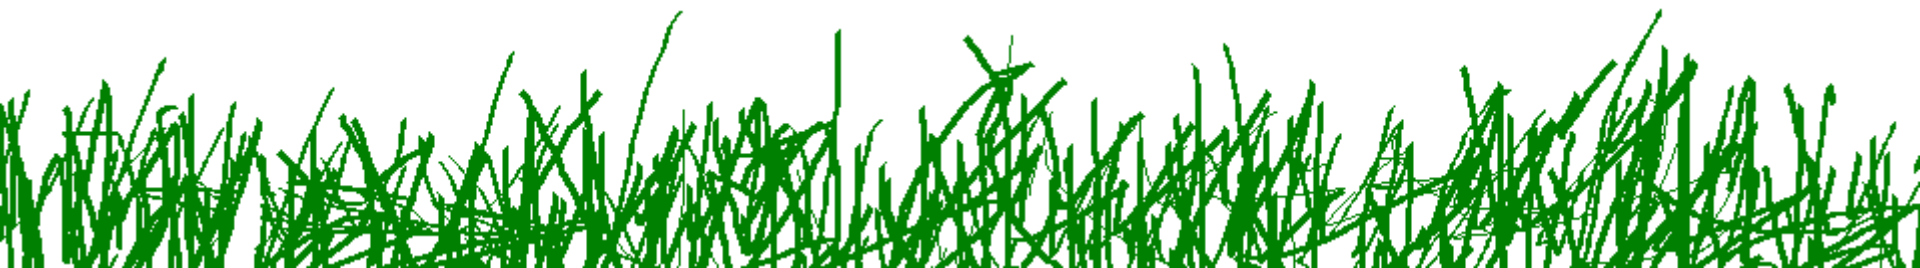
Drawbacks of Vector Clock

- Need to **maintain memory** units for tracking all events in all processes as a vector.
- **The total numbers of processes may not be known in advance.** Number of processes may be created or terminated at any time.
- Unnecessary wastage of maximum allocation of memory units.



Applications of Vector Clock

- **Distributed debugging.**
- Implementations of **causal ordering** communication and **causal distributed shared memory**.
- Establishment of **global breakpoints**.
- Determining the **consistency of checkpoints** in optimistic recovery.



Summary

- Physical Clocks
- Synchronizing Physical Clock (Algorithms)
- Problems with Physical Clock
- Lamport's Logical Clock
- Problems with Logical Clock
- Vector Clock
- Drawbacks of Vector Clocks
- Applications of Vector Clocks

References

1. Kshemkalyani, Ajay D., and Mukesh Singhal. Distributed computing: principles, algorithms, and systems. Cambridge University Press, 2011
2. Mukesh Singhal & N.G. Shivaratri, Advanced Concepts in Operating Systems
3. George Coulouris, Jean Dollimore and Tim Kindberg, “Distributed Systems Concepts and Design”, Fifth Edition, Pearson Education, 2012



Thank You

