

Logical Clocks (causality: as per chronology)

By Leslie Lamport

"happened before" relationship \rightarrow causally affects.

- If 'a' happened before 'b', then $\text{Clock}_a < \text{Clock}_b$.


$a \rightarrow b$

- If two processes interact with 'a' in one to 'b' in another, then

$\text{send}(m) \rightarrow \text{recv}(m)$

$a \rightarrow b$

$C_a < C_b$




- Rules for incrementing:

$C_i = C_i + d$ for processing events
($d = 1$)

$C_{j+1} = \max(C_j, tm) + 1$ for communication events.
 \downarrow
(message timestamp)

Partial Orders



• Transitive $e_{11} \rightarrow e_{22}$

• Antisymmetric $e_{11} \not\rightarrow e_{21}$

• Irreflexive $e_{11} \rightarrow e_{11}$

Cons: $\nexists a \rightarrow b ; C_a \leq C_b$.

(not strongly consistent) $\nexists C_a \leq C_b ; a \rightarrow b ?$ FALSE.

i.e. cannot determine chronological order given any two clock values.

Vector Clocks

- Rules for incrementing

$C_i[i] = C_i[i] + 1$ for processing events.

$\forall_k C_j[k] = \max(E_m[k], C_j[k] + 1)$

↓

Increment self clock
on receive.

For some clock values: t_a, t_b

$t_a == t_b \Rightarrow \forall_k t_a[k] == t_b[k]$

$t_a \neq t_b \Rightarrow \exists_k t_a[k] \neq t_b[k]$.


Cons:

Overlooks concurrent events.

Cannot determine which event occurred in which process, in what order.

Need more memory.

No. of processes may not be known in advance.



Applications :

- distributed debugging
- causal distributed shared memory
- global breakpoints
- consistency of checkpoints in optimistic recovery.

Vector clocks are strongly consistent.

i.e. if $t_a < t_b \Rightarrow a \rightarrow b$.

↓
solves the Lamport clock problem.