


MEMORY HANDLING

NVIDIA GPU Memory Structures

Memory hierarchy

- Thread:
 - Registers
- Thread:
 - Local memory
- Block of threads:
 - Shared memory
- All blocks:
 - Global memory



NVIDIA GPUs: Terminology

- Memory hardware
- Global Memory
 - DRAM available to all threads (SIMD processors in GPU)
- Local Memory
 - Private to the thread
- Shared Memory
 - Accessible to all threads of a Streaming Processor
- Thread Processor Registers

CACHES

- A cache is a high-speed memory that is physically close to the processor core.
- Caches are expensive.
- The maximum speed of a cache is proportional to the size of the cache.
- L1 cache is the fastest, but is limited in size to usually around 16 K, 32 K, or 64 K. It is usually allocated to a single CPU core.
- The L2 cache is slower, but much larger, typically 256 K to 512 K.
- The L3 cache may or may not be present and is often several megabytes in size.
- The L2 and/or L3 cache may be shared between processor cores or maintained as separate caches linked directly to given processor cores.
- Generally, at least the L3 cache is a shared cache between processor cores on a conventional CPU.
- This allows for fast intercore communication via this shared memory within the device



CACHES

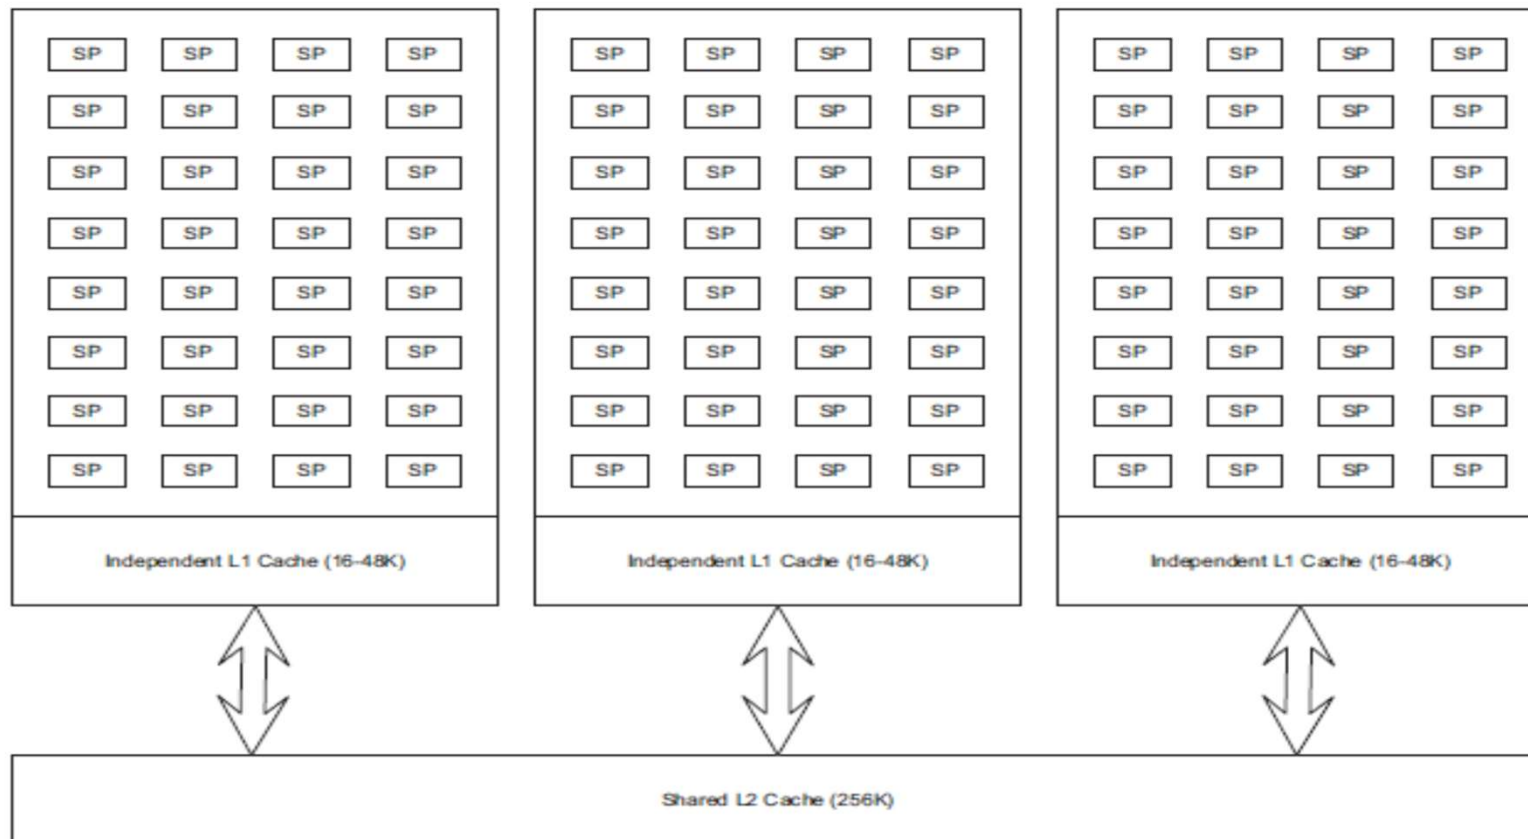


Fig: SM L1/L2 data path.

CACHES

- In Fermi GPU we introduce the concept of a nonprogrammer managed data cache.
- The architecture additionally has, per SM, an L1 cache that is both programmer managed and hardware managed.
- It also has a shared L2 cache across all SMs

CACHES

- **Types of data storage**
- On a GPU, we have a number of levels of areas where you can place data, each defined by its potential bandwidth and latency, as shown in Table below.

CACHES

Table 6.1 Access Time by Memory Type

Storage Type	Registers	Shared Memory	Texture Memory	Constant Memory	Global Memory
Bandwidth	~8 TB/s	~1.5 TB/s	~200 MB/s	~200 MB/s	~200 MB/s
Latency	1 cycle	1 to 32 cycles	~400 to 600	~400 to 600	~400 to 600

REGISTER USAGE

- The GPU, has thousands of registers per SM (streaming multiprocessor).
- An SM can be thought of like a multithreaded CPU core.
- On a typical CPU we have two, four, six, or eight cores.
- On a GPU we have N SM cores.
- On a Fermi GF100 series, there are 16 SMs on the top-end device. The GT200 series has up to 32 SMs per device
- A typical CPU will support one or two hardware threads per core.
- A GPU by contrast has between 8 and 192 SPs per core, meaning each SM can at any time be executing this number of concurrent hardware threads.

REGISTER USAGE

- On GPUs, application threads are pipelined, context switched, and dispatched to multiple SMs.
- The number of active threads across all SMs in a GPU device is usually in the tens of thousands range.

REGISTER USAGE

- One major difference we see between CPU and GPU architectures is how CPUs and GPUs map registers:
 - The CPU runs lots of threads by using register renaming and the stack.
 - To run a new task the CPU needs to do a context switch, which involves storing the state of all registers onto the stack (the system memory) and then restoring the state from the last run of the new thread.
 - This can take several hundred CPU cycles.
 - If you load too many threads onto a CPU it will spend all of the time simply swapping out and in registers as it context switches.
 - The effective throughput of useful work rapidly drops off as soon as you load too many threads onto a CPU.

REGISTER USAGE

- The GPU by contrast is the exact opposite.
 - It uses threads to hide memory fetch and instruction execution latency
 - The GPU does not use register renaming, but instead dedicates real registers to each and every thread.
 - Thus, when a context switch is required, it has near zero overhead.
 - On a context switch the selector (or pointer) to the current register set is updated to point to the register set of the next warp that will execute.

REGISTER USAGE

- We use registers to avoid usage of the slower memory types.
- For example, suppose we had a loop that set each bit in turn, depending on the value of some Boolean variable.

```
for (i=0; i<31; i++)  
{  
    packed_result |= (pack_array[i] << i);  
}
```

- Here we are reading array element *i* from an array of elements to pack into an integer, packed_result.
- We're **left shifting** the Boolean by the necessary number of bits and then using a **bitwise or** operation with the previous result.

REGISTER USAGE

- If the parameter ***packed_result*** exists in memory, you have **32 memory read and writes**.
- We can place the parameter ***packed_result*** in a local variable
- The compiler would place into a register.
- As we accumulate into the **register** instead of in main memory, and later write only the result to main memory, **we save 31 of the 32 memory reads and writes**.

REGISTER USAGE

- Assume 500 cycles for one global memory read or write operation.
- For every value you'd need to read, apply the or operation, and write the result back.
- You have **32 X read + 32 X write = 64X500 cycles =32,000 cycles.**
- The register version would eliminate 31 read and 32 write operations, replacing the 500-cycle operations with single-cycle operations.

$(1X \text{ memory read}) + (1 X \text{ memory write}) + (31x \text{ register read}) + (31X \text{ register write})$

or

$(1x500) + (1X500) + (31 X1) + (31X 1) = 1062 \text{ cycles versus } 32,000 \text{ cycles.}$

- This is a huge reduction in the number of cycles.
- We have a 31 times improvement to perform a relatively common operation in certain problem domains

SHARED MEMORY

- Shared memory is effectively a user-controlled L1 cache.
- The L1 cache and shared memory share a 64 K memory segment per SM.
- In Kepler this can be configured in 16 K blocks in favour of the L1 or shared Memory.
- In Fermi the choice is 16K or 48K in favour of the L1 or shared memory.
- The shared memory has in the order of 1.5 TB/s bandwidth with extremely low latency.
- This is hugely superior to the up to 190GB/s available from global memory, but around one-fifth of the speed of registers.

SHARED MEMORY

- Shared memory is a bank-switched architecture.
- On Fermi it is 32 banks wide, and on G200 and G80 hardware it is 16 banks wide.
- Each bank of data is 4 bytes in size, enough for a single-precision floating-point data item or a standard 32-bit integer value.
- Kepler also introduces a special 64 bit wide mode .
- There is no need for a one-to-one sequential access, just that every thread accesses a separate bank in the shared memory.
- There is a crossbar switch connecting any single bank to any single thread.
- This is very useful when you need to swap the words.



SHARED MEMORY

- Every thread in a warp reads the same bank address.
- This triggers a broadcast mechanism to all threads within the warp.
- Usually thread zero writes the value to communicate a common value with the other threads in the warp (fig below).
- If we have any other pattern, we end up with bank conflicts of varying degrees.
- This means you stall the other threads in the warp that idle while the threads accessing the shared memory address queue up one after another.
- One important aspect of this is that it is not hidden by a switch to another warp, so we do in fact stall the SM.