

SEARCHING

GLOBAL MEMORY: Searching

- The search we have two options: a binary search.
- A binary search takes advantage of the fact we have a sorted list of samples from the previous step.
- It works by dividing the list into two halves and asking whether the value it seeks is in the top or bottom half of the dataset.
- It then divides the list again and again until such time as it finds the value.
- The worst case sort time for a binary search is $\log_2(N)$.

Parallel Searching

```
#include <iostream>
```

```
#include <cstdio>
```

```
// Kernel function for binary search
```

```
__global__ void binarySearchKernel(int *arr, int target, int left, int right,  
int *result) {
```

```
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
```

Parallel Searching

```
while (left <= right) {  
    int mid = left + (right - left) / 2;  
    if (arr[mid] == target) {  
        atomicExch(result, mid); // Store the result in a thread-safe manner  
        return;  
    } else if (arr[mid] < target) {  
        left = mid + 1;  
    } else {  
        right = mid - 1;  
    }  
}
```

Parallel Searching

```
int main() {  
    const int arraySize = 1024;  
    const int target = 42;  
  
    int *hostArray, *deviceArray, *deviceResult;  
  
    // Allocate memory on host and device  
    hostArray = new int[arraySize];  
    cudaMalloc(&deviceArray, arraySize * sizeof(int));  
    cudaMalloc(&deviceResult, sizeof(int));  
}
```

Parallel Searching

```
// Initialize the sorted array on the host
```

```
for (int i = 0; i < arraySize; ++i) {
```

```
    hostArray[i] = i * 2;
```

```
}
```

```
// Copy data from host to device
```

```
    cudaMemcpy(deviceArray, hostArray, arraySize * sizeof(int),  
cudaMemcpyHostToDevice);
```

```
// Set up grid and block dimensions
```

```
int threadsPerBlock = 256;
```

```
int blocksPerGrid = (arraySize + threadsPerBlock - 1) /  
threadsPerBlock;
```

Parallel Searching

```
// Launch kernel
    binarySearchKernel<<<blocksPerGrid, threadsPerBlock>>>(deviceArray, target, 0,
arraySize - 1, deviceResult);
// Copy result from device to host
    int result;
    cudaMemcpy(&result, deviceResult, sizeof(int), cudaMemcpyDeviceToHost);
if (result != -1) {
    std::cout << "Element " << target << " found at index " << result <<
std::endl;
} else {
    std::cout << "Element " << target << " not found in the array." <<
std::endl;
}
```

Parallel Searching

```
// Clean up  
delete[] hostArray;  
cudaFree(deviceArray);  
cudaFree(deviceResult);  
return 0;  
}
```